

# Computer Science Project Proposal

## Path ORAM on MirageOS

R. Horlick, Homerton College

Originator: Dr N. Sultana

15 October 2015

**Project Supervisors:** Dr N. Sultana & Dr R. M. Mortier

**Director of Studies:** Dr B. Roman

**Project Overseers:** Dr M. G. Kuhn & Prof P. E. Sewell




Figure 1: The MirageOS stack with and without ORAM

## Introduction and Description of the Work

As the cost of large-scale cloud storage decreases and the rate of data production grows, individuals and small businesses will find themselves increasing reliant on the trust of cloud providers. We can, of course, use encryption and be safe in the knowledge that no adversary will be able to view the plaintext of our data, but this is not enough. It turns out that the pattern of access to the data can leak large amounts of information. In a study on an encrypted email repository, up to 80% of search queries could be inferred from the access pattern alone! So clearly this is a leak worth plugging, but how can we do it?

The solution to our problem is Oblivious Random Access Memory (ORAM), a cryptographic primitive that ensures that an adversary has negligible probability of learning anything about the logical access pattern, even with full access to the physical one. We normally talk about an ORAM as a block device, with  $N$  blocks of  $B$  bits each, giving an ORAM of size  $N \cdot B$ . We are most interested in the asymptotic bandwidth cost of the ORAM protocol, because we are operating over the internet.

A trivial ORAM algorithm operates by scanning over the whole ORAM and reading/updating only the relevant block, but this has  $O(N)$  bandwidth cost, which is highly impractical for large-scale storage. Luckily, much better algorithms have been proposed. We choose to focus on Path ORAM, because it has only  $O(\log N)$  bandwidth cost in the worst case if  $B = \Omega(\log^2 N)$ , as well as being incredibly simple conceptually.

ORAM protocols generally operate in a client-server model. For our purposes the server is the cloud storage provider and the client will be a MirageOS cloud instance. MirageOS is a unikernel operating system designed to run on the Xen hypervisor. I have chosen to use Mirage for a number of reasons. Firstly, it is lightweight and designed to be run in the cloud, meaning that simple cloud services can be built on top of it that fully leverage the ORAM. Secondly, it is written in OCaml, meaning that I can take full advantage of the rich module system. This will allow me to write my implementation as a functor that takes an implementation of Mirage’s BLOCK interface and creates a new BLOCK implementation that uses ORAM, as shown in Figure 1. This means that the ORAM could be used with any underlying implementation of the BLOCK interface and that it would plug seamlessly into any existing program.

# Starting Point

MirageOS is a framework, that pulls together a number of libraries and syntax extensions, to provide a lightweight unikernel operating system. It provides a command line tool for generating the main file, that links together implementations of various parts of the system, and passes them to the unikernel. There are a number of module signatures that define the operation of devices, such as `CONSOLE` for consoles, `ETHIF` for ethernet, and most importantly for us `BLOCK`, for block devices.

There are currently two implementations of the `BLOCK` interface, one for Unix and one for Xen, and I would like to support both. Thus I will build a functor that takes one of these and emulates a block device using it, adding ORAM to it. This will most likely use some of the algorithms/code from the existing implementations, for dealing with buffers, implemented using the `Cstruct` library.

## Substance and Structure of the Project

### Substance

The Path ORAM protocol has three main components: a binary tree, a stash and a position map. The binary tree is the main storage space. Every node in the tree is a bucket, which can contain up to  $Z$  blocks. The tree has height  $L$ , where the tree of height 0 consists only of the root node, and the leaves are at level  $L$ . The stash is temporary client-side storage, consisting only of a set of blocks waiting to be put back into the tree. The position map associates, with each block ID, an integer between 0 and  $2^L - 1$ . The invariant that the Path ORAM algorithm maintains is that if the position of a block  $x$  is  $p$ , then  $x$  is either in some bucket along the path from the root node to the  $p^{th}$  leaf, or in the stash. On every access to the tree, a whole path is read into, the accessed block is assigned a new random position and then as many blocks as possible are written back into the same path. This means that in any two access to the same block, the paths that are read are statistically independent.

We can extend the basic path ORAM algorithm with recursion. That is, calling the data ORAM  $ORAM_0$ , we store the position map of  $ORAM_0$  in a smaller ORAM,  $ORAM_1$ , and the position map for this in an even smaller ORAM,  $ORAM_2$ . We can do this until we have a sufficiently small position map on the client. Supposing that we store  $\chi$  leaf addresses in each PosMap ORAM, the position for a data block with address  $a_0$  is at  $a_1 = a_0/\chi$  in  $ORAM_1$ , and in general  $a_n = a_0/\chi^n$  for the address in  $ORAM_n$ .

Recursion does add an overhead, but we can reduce this overhead by exploiting locality. Assuming that programs will access adjacent data blocks, we can cache PosMap blocks in a PosMap Lookaside Buffer, so that if all  $\chi$  data blocks that are referenced in a PosMap block are accessed in turn we only need to do the recursion once. Doing this naïvely, however, breaks security, because we are revealing information through the cache hit pattern. To avoid this we use Unified ORAM, which combines all of the recursive ORAMs into a single logical tree. We then use the address space to separate the levels of recursion, so addresses 1 to  $N$  are for data blocks,  $N + 1$  to  $N + (N/\chi)$  for  $ORAM_1$  and so on. Now all accesses occur in the same tree, and the security of Path ORAM keeps the cache miss pattern hidden.

A final optimisation compresses the PosMaps, reducing the number of levels of recursion required to achieve the desired client side storage, resulting in an asymptotic

bandwidth complexity for ORAM with small block size.

The last two optimisations were originally designed and tested in a secure processor setting, so their application to the cloud storage setting is novel. The evaluation section of this project will provide empirical evidence as to whether these optimisations are worthwhile for our purposes.

## Structure

The project breaks down into the following sub-projects:

1. Familiarising myself with OCaml, MirageOS and related libraries
2. Implementing the basic Path ORAM functor and testing that it works in place of existing BLOCK device implementations
3. Implementing the three main optimisations to Path ORAM as extensions to the same functor, allowing for the use of different combinations of optimisations
4. Creation of a suite of tests and experiments to evaluate the relative performance of the optimisations and their conformity to theoretical bounds
5. Writing the dissertation

## Success Criterion for the Main Result

To demonstrate, through well chosen examples, that I have implemented a secure Path ORAM functor, along with a number of “optimisations”.

## Possible Extensions

Path ORAM (and other tree based ORAMs) are limited in the fact that they have fixed-sized trees. Thus, we either need to know our storage requirements before setting up the ORAM, potentially wasting resources, or resize them in the naïve way as storage requirements increase. If I achieve the goals of my main project, including evaluation, ahead of schedule I will examine the possibility of making Path ORAM resizable.

## Timetable: Work plan and Milestones

Planned starting date is 16/10/2015.

1. **Michaelmas weeks 2–4** Familiarise myself with relevant Mirage libraries. Implement basic Path ORAM functor.
2. **Michaelmas weeks 5–6** Implement basic test harness. Start implementation of recursive Path ORAM.
3. **Michaelmas weeks 7–8** Design comparative tests for first two version of the functor.

4. **Michaelmas vacation** Implement Unified ORAM and PosMap compression. Write implementation section of dissertation.
5. **Lent weeks 0–2** Write progress report. Design further comparative tests for set of optimisations.
6. **Lent weeks 3–5** Run main experiments and achieve working project.
7. **Lent weeks 6–8** Write main chapters of dissertation.
8. **Easter vacation:** Finish first draft of dissertation and submit to supervisor for feedback.
9. **Easter term 0–3:** Proof reading and then an early submission so as to concentrate on examination revision.

## Resources Required

- My own laptop for implementation and testing
- My own external hard disk for backups
- GitHub for version control and backup storage
- MirageOS libraries as a basis for the project